

YOLObile: Real-Time Object Detection on Mobile Devices via Compression-Compilation Co-Design

Yuxuan Cai^{1*}, Hongjia Li^{1*}, Geng Yuan^{1*}, Wei Niu², Yanyu Li¹, Xulong Tang³, Bin Ren², Yanzhi Wang¹

¹Northeastern University, ²William & Mary, ³University of Pittsburgh

Presented by Rui Chen

AAAI 2021

Outline

- **Motivation**
- Framework Overview
 - Block-punched Pruning
 - Compiler-assisted Acceleration
 - Mobile GPU-CPU Collaborative Scheme
- Experimental Results

Background: Object Detection

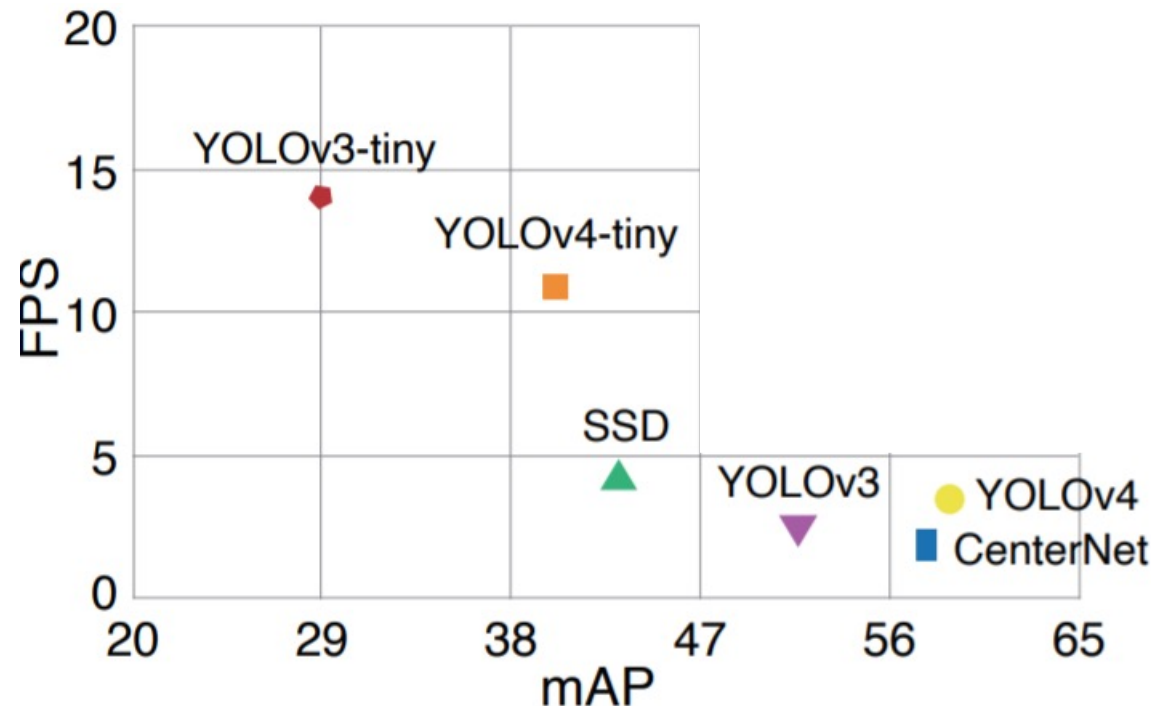
- Autonomous Driving
- UAV Obstacle Avoidance
- Augmented Reality
- Robot Vision



Background: Object Detection

FPS: frames per second
mAP: mean average precision

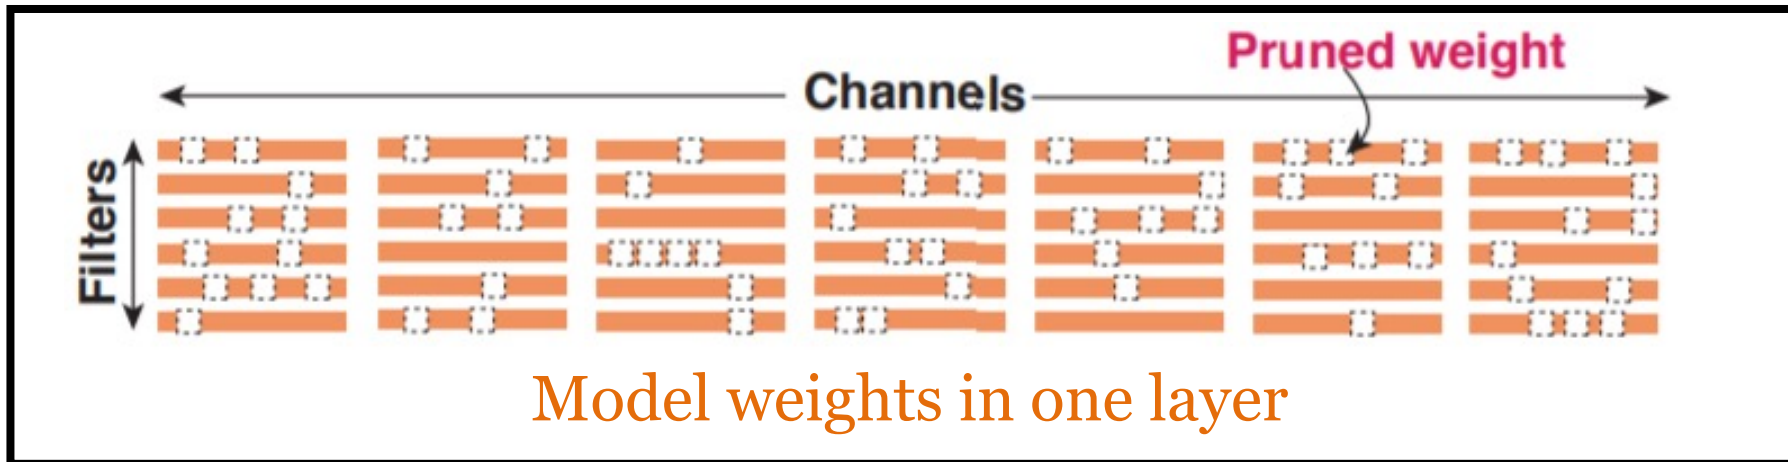
- Executing DNNs inference on mobile device is still challenging
 - High computation and storage demands
 - Cannot achieve real-time performance with high accuracy



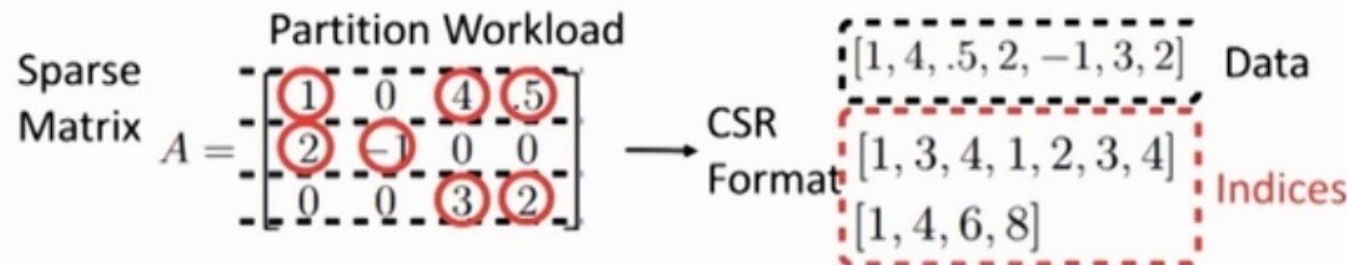
FPS VS mAP on MS COCO dataset

Background: Pruning

- **Non-structured** weight pruning: arbitrary weight can be pruned



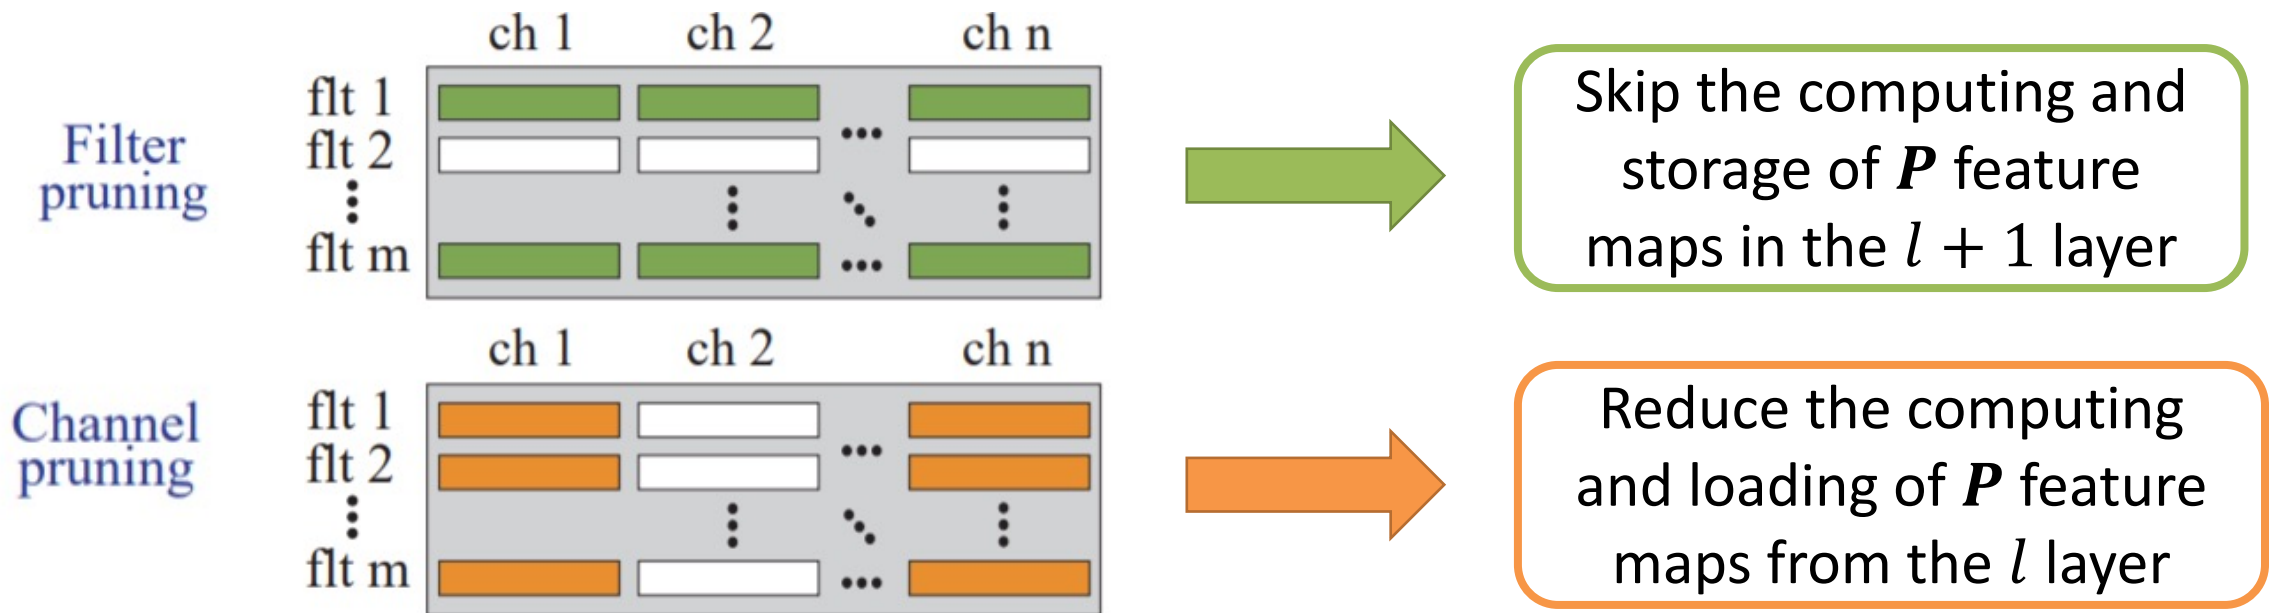
- **Limited** actual development
 - Indices are required for sparse format – **speed degradation** in GPU/CPU



Background: Pruning

- **Structured** weight pruning
 - Hardware friendly
 - **Large Accuracy loss** due to coarse granularity

$$I * W = O$$
$$I \in R^{a \times b \times n}, W \in R^{k \times k \times n \times m}, O \in R^{c \times d \times m}$$



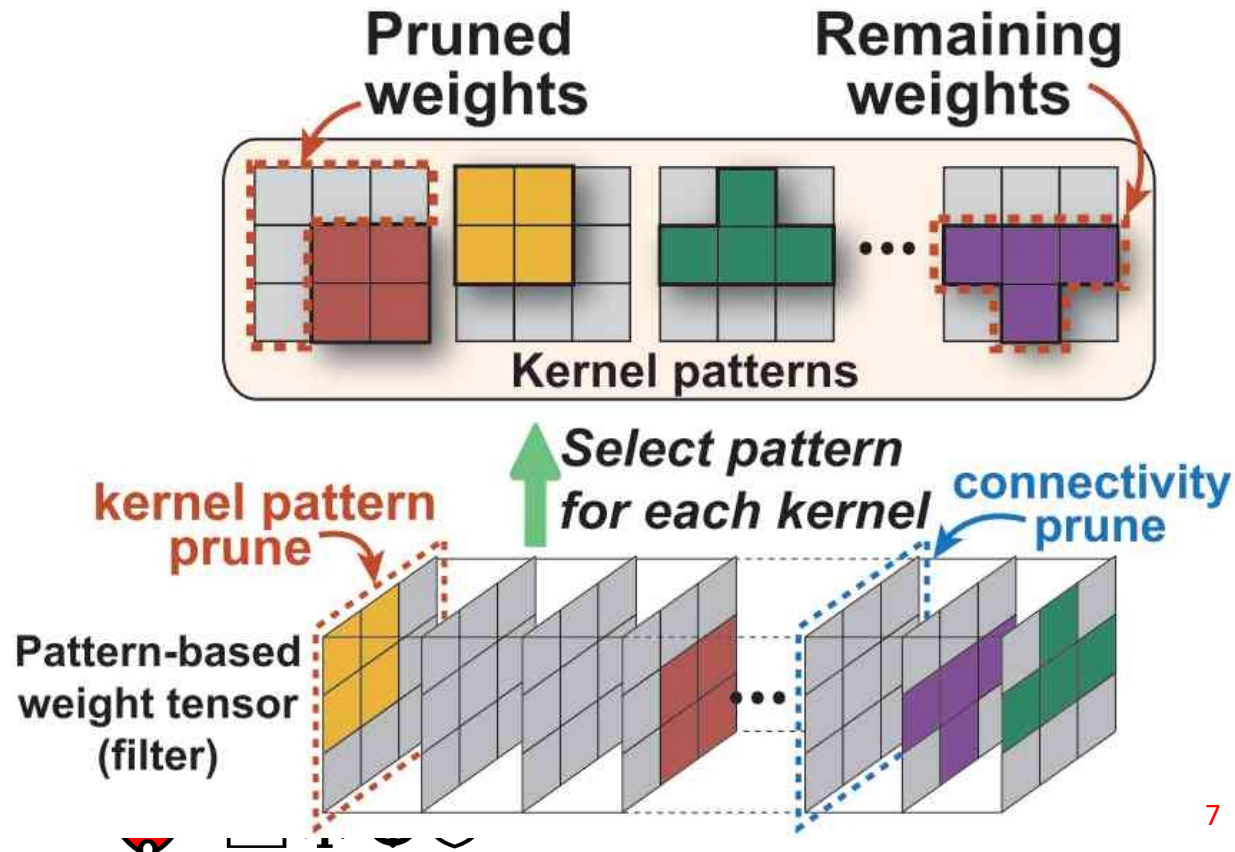
Ch: channel
Flt: filter

Background: Pruning

mAP: mean average precision

- **Pattern-based** weight pruning¹

It is inspired by Sparse Convolution Pattern (using Gaussian and Laplace filters).



Pro: Accurate and efficient

Cons: Only applicable to 3x3 conv layers

YOLOv4:

21% mAP loss with 5x compression, not acceptable

¹PCONV: Ma et al. "The Missing but Desirable Sparsity in DNN Weight Pruning for Real-Time Execution on Mobile Devices", AAAI, 2020

Contribution

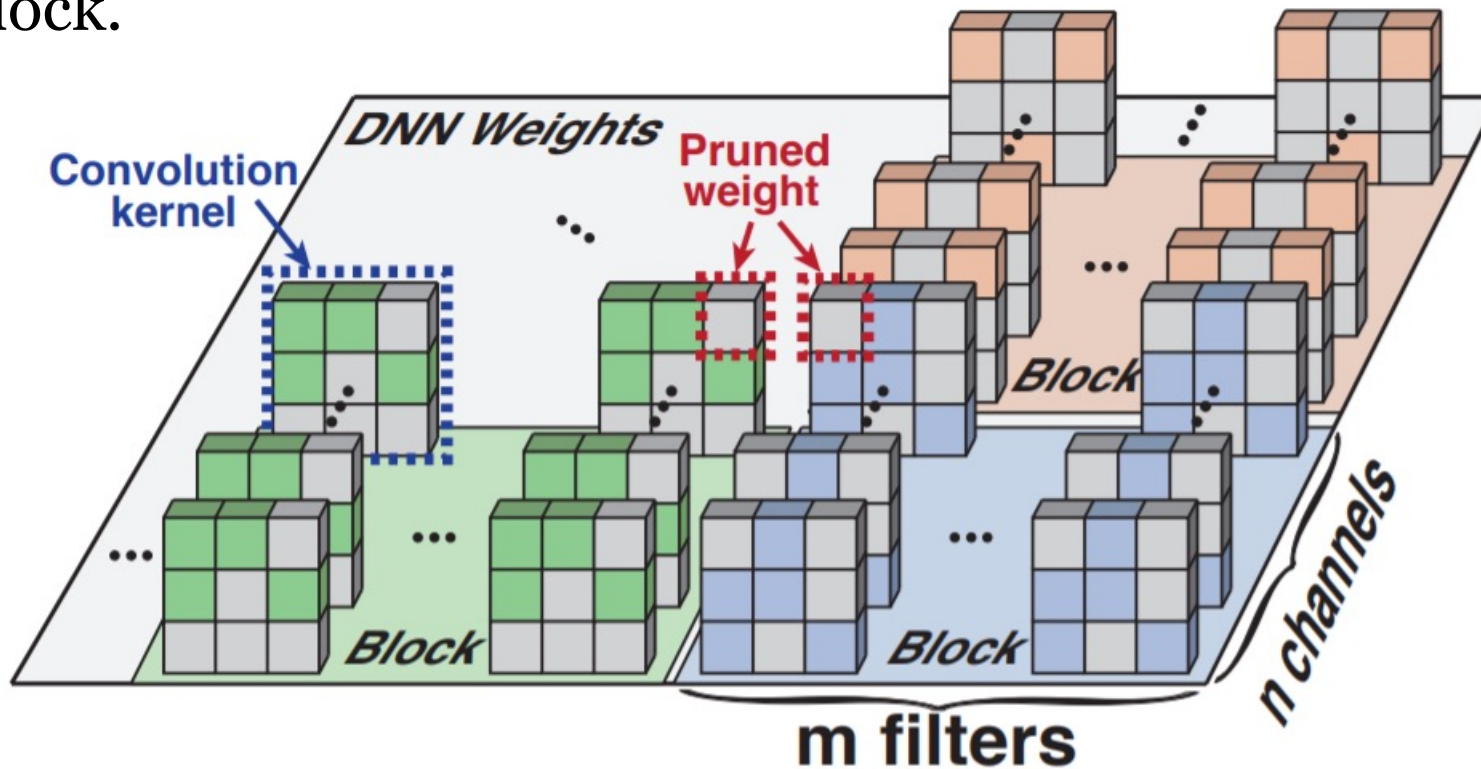
- Propose “Block-Punched Pruning” that can be applied on different types of layers
- Propose an efficient computation method for further accelerating the DNN inference speed of object detection tasks.

Outline

- Motivation
- **Framework Overview**
 - Block-punched Pruning
 - Compiler-assisted Acceleration
 - Mobile GPU-CPU Collaborative Scheme
- Experimental Results

Block-Punched Pruning

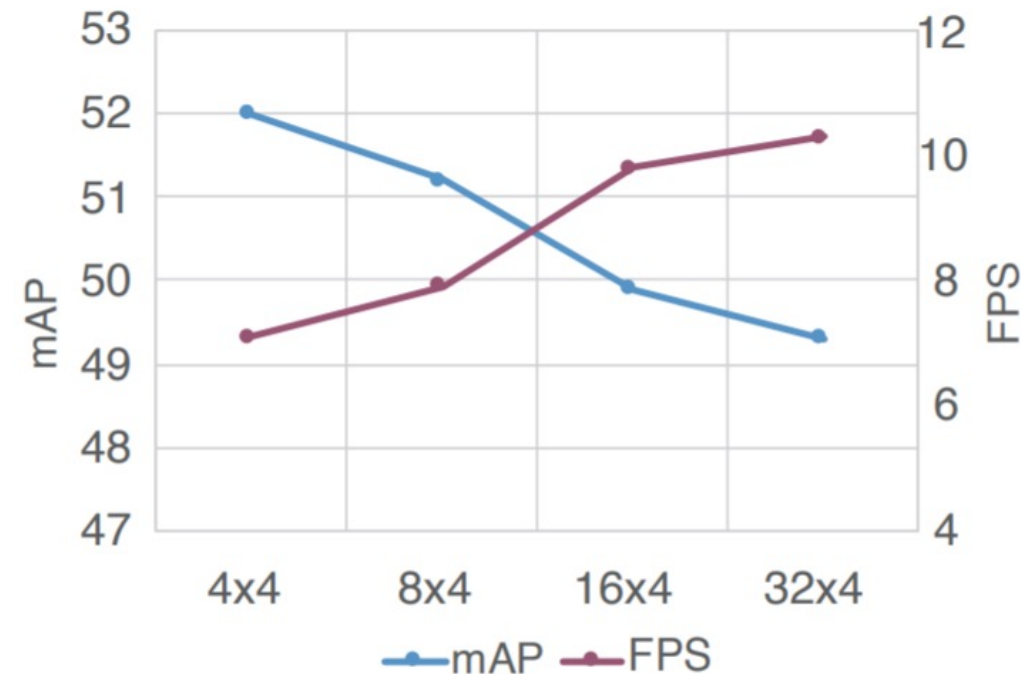
- Divided the 4d weight matrix into blocks
- The weights to be pruned will punch through the same location of all kernels within a block.



Block-Punched Pruning

- Block size?
 - #channels: equals to the length of the vector registers of each treading in the mobile CPU/GPU
 - #filters: the trade-off between accuracy and hardware acceleration (multithreading)
- Larger block size, better leverage the hardware parallelism, more accuracy loss

FPS and mAP in different block size



Reweighted Regularization Pruning

- How to prune?

Controls the trade-off between accuracy and sparsity

Pruning problem:

$$\underset{\mathbf{W}, \mathbf{b}}{\text{minimize}} \quad f(\mathbf{W}; \mathbf{b}) + \lambda \sum_{i=1}^N R(\alpha_i^{(t)}, \mathbf{W}_i),$$

Regularization term

- Reweighted group lasso¹ regularization $\mathbf{W}_i = [\mathbf{W}_{i1}, \mathbf{W}_{i2}, \dots, \mathbf{W}_{iK}]$, $\mathbf{W}_{ij} \in \mathbb{R}^{g_i m \times g_i n}$.

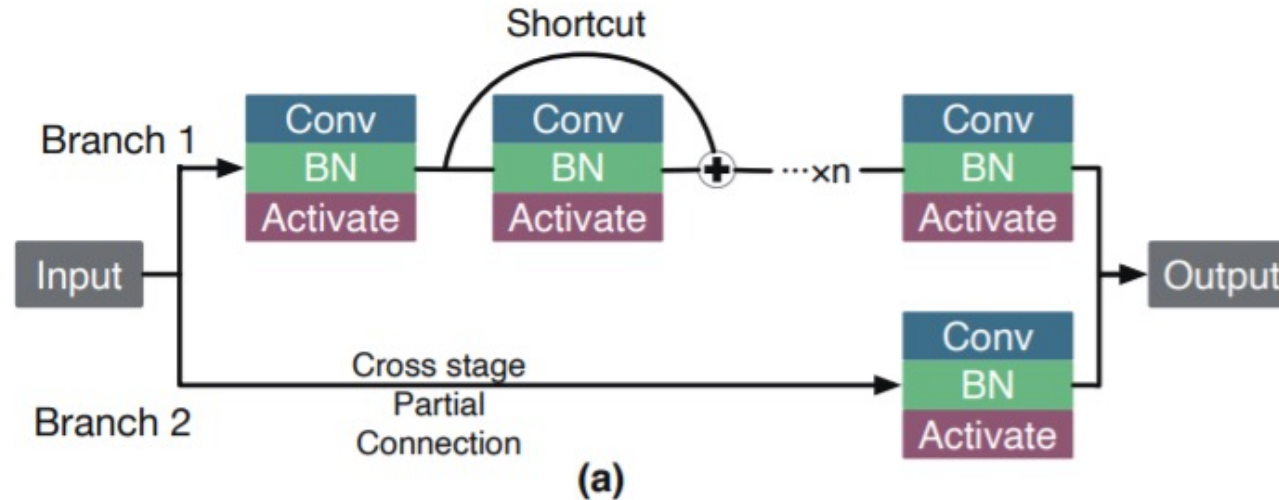
$$R(\alpha_i^{(t)}, \mathbf{W}_i) = \sum_{j=1}^K \sum_{h=1}^{g_m^i} \sum_{w=1}^{g_n^i} \|\alpha_{ijn}^{(t)} \circ [\mathbf{W}_{ij}]_{h,w}\|_F^2,$$

Penalize the weight with small values

$$\alpha_{ijn}^{(t)} = \frac{1}{\|[\mathbf{W}_{ij}]_{h,w}^{t-1}\|_F^2 + \epsilon}.$$

- Use pre-trained model and retrain for 3-4 iterations

Mobile CPU-GPU Collaborative Scheme



Yolov4

- Off-line device selection based on latency

$$\begin{cases} T_{par} = \max\{t_{g1}, t_{c2} + \tau\} \\ T_{ser} = t_{g1} + t_{g2} \end{cases}$$



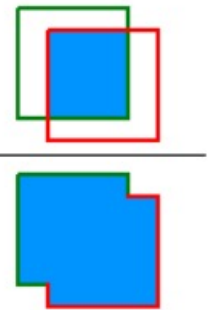
select the optimal executing device for Branch 2

Outline

- Motivation
- Framework Overview
 - Block-punched Pruning
 - Compiler-assisted Acceleration
 - Mobile GPU-CPU Collaborative Scheme
- **Experimental Results**

Experiment Setup

- Nvidia Docker+Pytorch
- Training cost: 4xRTX2080Ti, 5 days
- Test device: Samsung Galaxy S20
- Manually designed layerwise prune ratio
- MS COCO dataset: 320x320 input image
- mAP: mean Average Precision under Intersection Over Union (IoU) 0.5 for multi-labels.
- AP[0.5:0.95]: Average precision under IoU 0.5 to 0.95

$$IOU = \frac{\text{area of overlap}}{\text{area of union}} = \frac{\text{area of overlap}}{\text{area of union}}$$


Performance evaluation

Approach	Input Size	backbone	#Weights	#FLOPs	mAP	AP@[.5:.95]	FPS
CenterNet-DLA (Duan et al.)	512	DLA34	16.9M	52.58G	57.1	39.2	1.9
CornerNet-Squeeze (Law et al.)	511	-	31.77M	150.15G	-	34.4	0.3
SSD (Liu et al.)	300	VGG16	26.29M	62.8G	43.1	25.1	4.2
MobileNetv1-SSDLite (Sandler et al.)	300	MobileNetv1	4.31M	2.30G	-	22.2	49
MobileNetv2-SSDLite (Sandler et al.)	300	MobileNetv2	3.38M	1.36G	-	22.1	41
Tiny-DSOD (Li et al.)	300	-	1.15M	1.12G	40.4	23.2	-
YOLOv4 (Bochkovskiy, Wang, and Liao)	320	CSPDarknet53	64.36M	35.5G	57.3	38.2	3.5
YOLO-Lite (Huang, Pedoeem, and Chen)	224	-	0.6M	1.0G	-	12.26	36
YOLOv3-tiny (Redmon and Farhadi)	320	Tiny Darknet	8.85M	3.3G	29	14	14
YOLOv4-tiny (Bochkovskiy, Wang, and Liao)	320	Tiny Darknet	6.06M	4.11G	40.2	-	11
YOLOBile (GPU only)	320	CSPDarknet53	4.59M	3.95G	49	31.6	17
YOLOBile (GPU&CPU)	320	CSPDarknet53	4.59M	3.95G	49	31.6	19.1

Table 2: Accuracy (mAP) and speed (FPS) comparison with other object detection approaches.

Performance evaluation

- All the results are evaluated under our compiler optimization

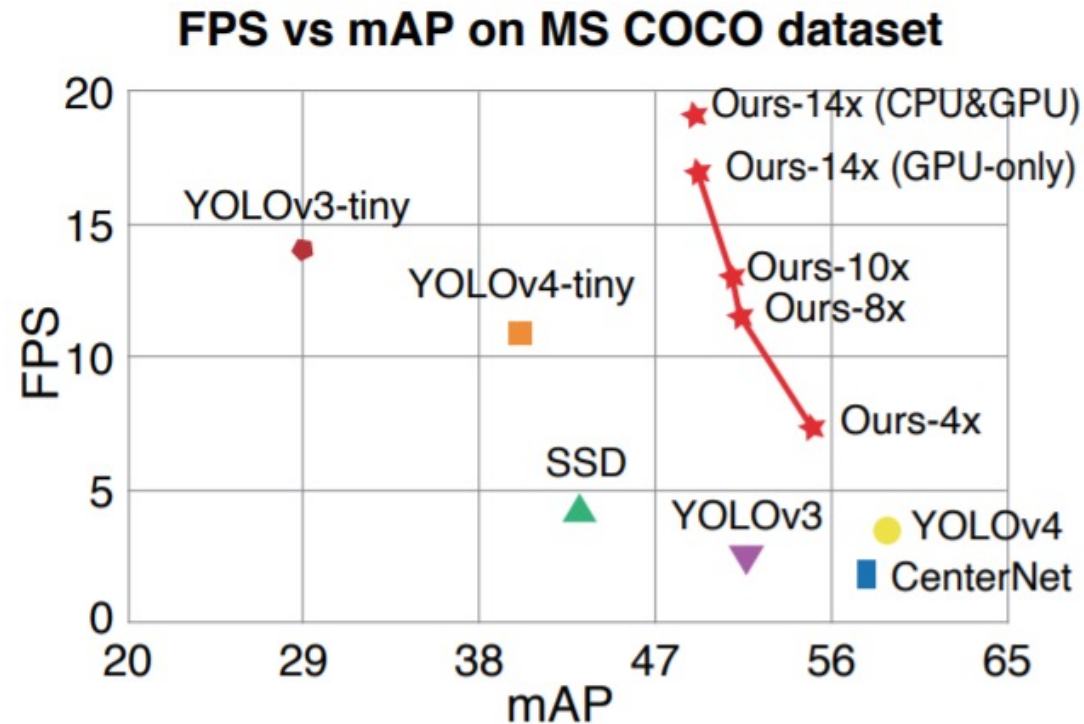


Figure 6: The accuracy (mAP) and speed (FPS) comparison of YOLOmobile under different compression rate and different approaches.

Performance evaluation

#Weights	#Weights Comp. Rate	#FLOPs	mAP	AP@[.5:.95]	FPS
64.36M	1×	35.8G	57.3	38.2	3.5
16.11M	3.99×	10.48G	55.1	36.5	7.3
8.04M	8.09×	6.33G	51.4	33.3	11.5
6.37M	10.1×	5.48G	50.9	32.8	13
4.59M	14.02×	3.95G	49	31.9	17

Table 1: Accuracy and speed under different compression rates.

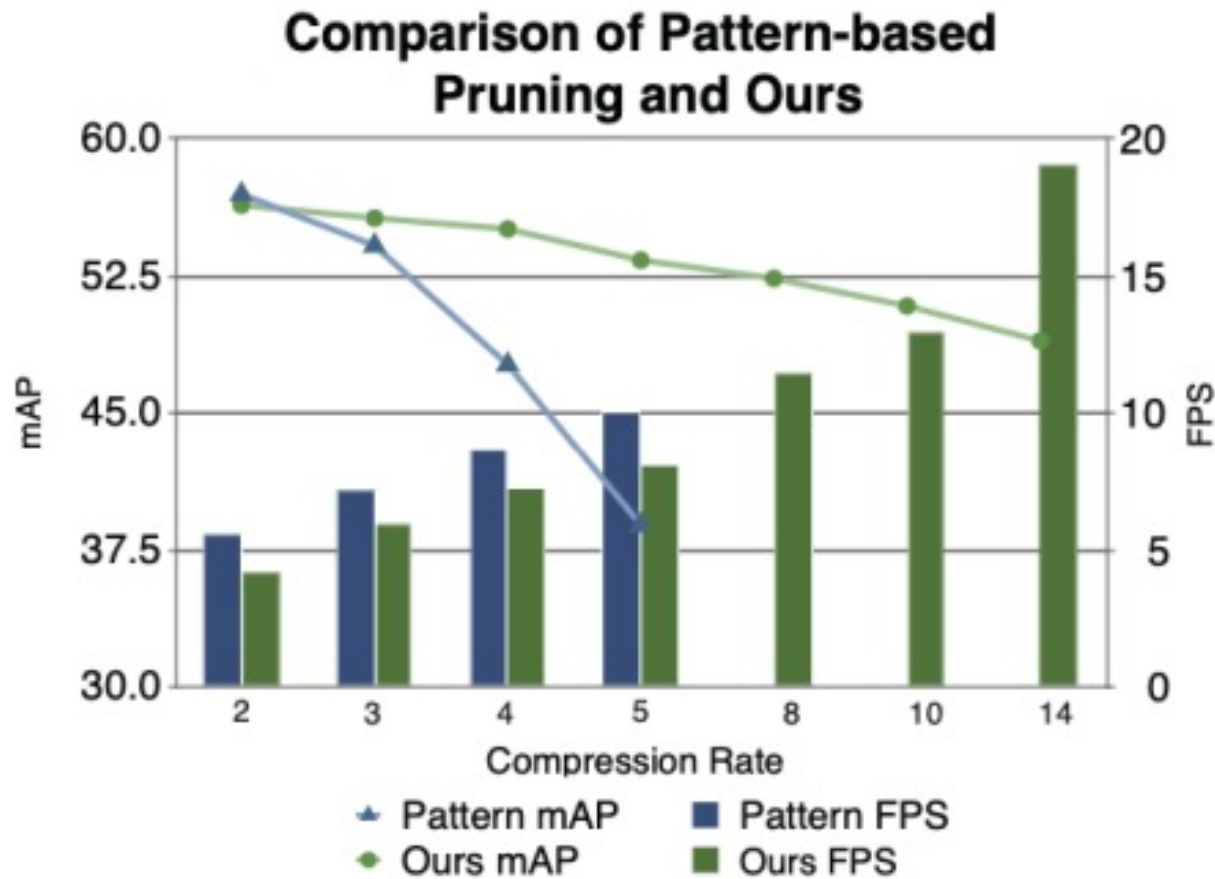
Ablation Study

Pruning Scheme	#Weights	#Weights Comp. Rate	mAP	FPS
Not Prune	64.36M	1×	57.3	3.5
Unstructured	8.04M	8.09×	53.9	6.4
Structured	8.04M	8.09×	38.6	12
Ours	8.04M	8.09×	51.4	11.5

Table 3: Comparison of different pruning schemes.

Compared with structured pruning and unstructured pruning, our block-punched pruning scheme achieves both high accuracy and fast inference speed.

Ablation Study



- Block-punched method achieves a small accuracy drop when the compression rate is high (>5).
- When the compression rate is 5, the Pattern-based method results in a sharp drop down of the curve.

Conclusion

- YOLObile, a real-time object detection framework on mobile devices via compression-compilation co-design
 - Block-Punched Pruning
 - Reweighted Regularization Pruning Algorithm
 - Compiler-assisted Acceleration
 - Mobile GPU-CPU Collaborative scheme

An aerial photograph of the University of Houston campus at dusk. The foreground shows several large, modern university buildings with flat roofs and some with glass facades. A central green lawn with winding paths is visible. In the background, the Houston city skyline is visible against a twilight sky with soft orange and blue hues. A large, solid red rectangular banner is positioned at the top of the image, containing the text "THANK YOU" in white, bold, sans-serif capital letters.

THANK YOU

UNIVERSITY of **HOUSTON** | ENGINEERING